

00 - Introduction

François Pitié

Ussher Assistant Professor in Media Signal Processing
Department of Electronic & Electrical Engineering, Trinity College Dublin

[4C16] Machine Learning with Applications in Media Engineering


This module is an introduction to Machine Learning and especially **Deep Neural Nets**, which are disrupting all aspects of society today.

The material is constructed in collaboration with leading industrial practitioners including Google, YouTube and Movidius/Intel.

Hands on labs will give you experience with these applications.

Deep Learning is a particular type of **machine learning** method, and is thus part of the broader field of **artificial intelligence** (using computers to reason).

Deep learning is another name for **artificial neural networks**, which are a loosely inspired by the structure of the neurons in the cerebral cortex.



The recent quantum leap in machine learning has solely been driven by deep learning successes.

When you read or hear about AI or machine Learning successes in recent years, it really means Deep Learning successes.

Machine Learning can be split into 3 main fields:

Supervised Learning:

From a labelled dataset (\mathbf{x}_i, y_i) , we want to infer $f(\mathbf{x}_i) = y_i$

Unsupervised Learning:

What can we learn about a dataset (\mathbf{x}_i) by just looking at it? (ie. without any labelled information y_i)

Reinforcement Learning:

How can an agent interact with its environment (the data) to get maximum reward (eg. game playing, robots learning to walk).


Deep Learning has made major breakthroughs in all three fields.

Deep Learning Successes

Image Recognition is one of the core applications of Computer Vision. ImageNet [www.image-net.org] runs an annual challenge where software programs compete to correctly classify and detect objects and scenes in images.




The error rate in object recognition for that challenge has massively dropped since the introduction of deep neural networks in 2012 [1]. Machines can now do better than humans.




[1] ImageNet Classification with Deep Convolutional Neural Networks
A Krizhevsky, I Sutskever, G Hinton, 2012 [<https://goo.gl/wxn2Y>]

Neural nets are also advancing the state of the art in
Scene Understanding.



[1] Conditional Random Fields as Recurrent Neural Networks
S Zheng et al., 2015 [<https://arxiv.org/abs/1502.03240>]

Image models combined with language models make it possible to automatically **generate captions** from images.



All major tech companies have changed their machine translation systems to use Deep Learning.

Google used to average a yearly 0.4% improvement on their machine translation system. Their first attempt at using Deep Learning yielded an overnight 7% improvement! More than in an entire lifetime!

Several years of handcrafted development could not match a single initial deep learning implementation.

New York Times, “The Great AI Awakening” [<https://goo.gl/DPYp6d>]

In 2014, Skype Translator was announced. It trains and optimises **speech recognition**, automatic **machine translation** and **speech synthesis** tasks, acting as the glue that holds these elements together.

Skype demo [<https://www.youtube.com/embed/NhxKg2PA3ZI?start=0>]
Microsoft blog post [<https://go.gl/eAuPcS>]

Deep learning has also been introduced in reinforcement learning to solve complex sequential decision making problems.

Recent successes include:
playing old Atari computer games,
programming real world Robots
and beating humans at Go.



demo: Robots Learning how to walk https://www.youtube.com/hx_bgoTF7bs

DeepMind [<https://goo.gl/3TcCNA>]

Reasons of a Success

Neural Networks have been around for decades. But is only now that it surpasses all other machine learning techniques.


Deep Learning is now a **disruptive** technology that has been unexpectedly taking over operations of technology companies around the world.

“The revolution in deep nets has been very profound, it definitely surprised me, even though I was sitting right there.”

— Sergey Brin, Google co-founder

Why now?

Because Deep Learning does scale.



Neural Nets are the only ML technique whose performance **scales efficiently with the training data size**. Other ML popular techniques just can't scale that well.

The advent of **big databases**, combined with cheaper **computing power** (Graphic Cards), meant that Deep Learning could take advantage of all this, whilst other techniques stagnated. Instead of using thousands of observations, Deep Learning can take advantage of billions.

The tipping point was 2012 in Computer Vision and around 2014 in Machine Translation.

Deep Learning offers a (relatively) simple framework to define and parametrise pretty much any kind of numerical method and then optimise it over massive databases.

By adopting an automated brute force approach to tuning algorithms, Deep Learning is able to surpass hand-tailored algorithms of skilled researchers.

It offers a systematic approach when before algorithms took years of human efforts to design.

Reasons of a Success

Democratisation

Deep Learning is a (relatively) simple framework.

Good programmers train state of the art neural nets without having done 10+ years of research in the domain.

It is an opportunity for start-ups and it has become a ubiquitous tool in tech companies.

Reasons of a Success

Global Reach

It has been applied successfully to many fields of research, industry and society:

self-driving cars, image recognition, detecting cancer, speech recognition, speech synthesis, machine translation, drug discovery and toxicology, customer relationship management, recommendation systems, bioinformatics, advertising, controlling lasers, etc.

Impact

Here is a question for you:

How long before your future job gets replaced by an algorithm?

Probably much sooner than you think. You might feel safe if you are an artist...

... but then again:



automatic style transfer [1]

[1] A Neural Algorithm of Artistic Style

L. Gatys, A. Ecker, M. Bethge, 2015 [<https://arxiv.org/abs/1508.06576>]

Does an AI need to make love to Rembrandt's girlfriend to make art? [<https://goo.gl/gi7rWE>]

Intelligent Machines: AI art is taking on the experts [<https://goo.gl/2kfyXd>]

4C16: Course Structure

Course Content

Part 1. Machine Learning Fundamentals

In Week 1 - 4, we will cover

Least Squares: the root of all Machine Learning.

(feature mapping, over/under fitting, regularisation, maximum likelihood)

Logistic Regression: your first Neuron.

(linear classifier, cross-entropy , gradient descent optimisation)

Classic Classifiers: overview of classic machine learning algorithms.

(SVM, Decision Trees, Kernel Trick, Nearest-Neighbours)

Comparing Classifiers:

(ROC curves, confusion tables, F1 score)

Course Content

Part 2. Neural Net Fundamentals

Feedforward Neural Network

(network architecture, back-propagation, regularisation, vanishing gradients)

Convolutional Neural Network

(convolution layers, pooling, visualisation, knowledge transfer)

Autoencoders

(unsupervised learning, PCA, CNN autoencoder)

Recurrent Neural Network

(LSTM, encodings, text processing)

Assignments

Week 2	Least Squares lab	[2% mark]
Week 3	Logistic Regression lab	[2% mark]
Week 4	Classifiers lab	[2% mark]
Week 5	Feed Forward Neural Net lab	[2% mark]
Week 6	Convolutional Neural Net lab	[2% mark]
Week 8	Quiz	[5% mark]
Week 8-10	lab project	[5% mark]
Week 9-11	lab project	[5% mark]
End of term	Exam	[75% mark]

We have developed a fantastic environment specially for you, so that you can learn best industry practices.



Google Cloud Platform



Keras



git

Labs

You will be programming in python 3 using Keras and TensorFlow. Everything will be running on the Google Cloud Platform, which gives you on-demand scalable computing resources.

Your coding environment will be a combination of shell/terminal, editor, and Jupyter notebook.

You will use Git to checkpoint your progress and for continuous feedback on lab assignments.

Books & Resources

[1] Deep Learning (MIT press)

Ian Goodfellow et al.

[\[https://www.deeplearningbook.org\]](https://www.deeplearningbook.org)

[2] Machine Learning on Coursera

Andrew Ng

[\[https://www.coursera.org/learn/machine-learning\]](https://www.coursera.org/learn/machine-learning)

[3] Neural Networks and Deep Learning

Michael Nielsen

[\[http://neuralnetworksanddeeplearning.com/\]](http://neuralnetworksanddeeplearning.com/)

[4] Curated list of links <https://github.com/ChristosChristofidis/awesome-deep-learning>

[5] Brandon Rohrer's YT channel <https://www.youtube.com/user/BrandonRohrer>

[6] Siraj Raval's YT channel <https://www.youtube.com/channel/UCWN3xxRkmTPmbKwht9FuE5A>