

## 02 - Logistic Regression

---

François Pitié

Ussher Assistant Professor in Media Signal Processing  
Department of Electronic & Electrical Engineering, Trinity College Dublin

# Motivation

With **Linear Regression**, we looked at linear models, where the output of the problem was a **continuous** variable (eg. height, car price, temperature, ...).

Very often you need to design a **classifier** that can answer questions such as: what car type is it? is the person smiling? is a solar flare going to happen? In such problems the model depends on **categorical** variables.

**Logistic Regression** (David Cox, 1958), considers the case of a binary variable. That is, the outcome is 0/1 or true/false.

There is a whole zoo of classifiers out there. Why are we covering logistic regression in particular?

Because logistic regression is the building block of Neural Nets.

## Introductory Example


We'll start with an example from Wikipedia:

*A group of 20 students spend between 0 and 6 hours studying for an exam. How does the number of hours spent studying affect the probability that the student will pass the exam?*

# Introductory Example

The collected data looks like so:

```
Studying Hours      : 0.75 1.00 2.75 3.50 ...  
result (1=pass,0=fail) : 0    0    1    0    ...
```




# Linear Approximation

Although the output  $y$  is binary, we could still attempt to fit a linear model via least squares:

$$h_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^T \mathbf{w} = w_0 + w_1 x_1 + \dots$$

# Linear Approximation

This is what the least squares estimate  $h_{\mathbf{w}}(\mathbf{x})$  looks like:



The model prediction  $h_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^T \mathbf{w}$  is continuous, but we could apply a threshold to obtain the binary classifier as follows:

$$y = [\mathbf{x}^T \mathbf{w} > 0.5] = \begin{cases} 0 & \text{if } \mathbf{x}^T \mathbf{w} \leq 0.5 \\ 1 & \text{if } \mathbf{x}^T \mathbf{w} > 0.5 \end{cases}$$

and the output would be 0 or 1.

Obviously this is not optimal as we have optimised  $\mathbf{w}$  so that  $\mathbf{x}^T \mathbf{w}$  matches  $y$  and not so that  $[\mathbf{w}^T \mathbf{x} > 0.5]$  matches  $y$ .

Let's see how this can be done.



# General Linear Model


The general problem of **general linear models** can be presented as follows. We are trying to find a linear combination of the data  $\mathbf{x}^T \mathbf{w}$ , such that the sign of  $\mathbf{x}^T \mathbf{w}$  tells us about the outcome  $y$ :

$$y = [\mathbf{x}^T \mathbf{w} + \epsilon > 0]$$

The quantity  $\mathbf{x}^T \mathbf{w}$  is sometimes called the **risk score**. It is a scalar value. The larger the value of  $\mathbf{x}^T \mathbf{w}$  is, the more certain we are that  $y = 1$ .


The error term is represented by the random variable  $\epsilon$ . Multiple choices are possible for the distribution of  $\epsilon$ .

In **logistic** regression, the error  $\epsilon$  is assumed to follow a **logistic distribution** and the risk score  $\mathbf{x}^T \mathbf{y}$  is also called the **logit**.



**Figure:** pdf of the logistic distribution

In **probit** regression, the error  $\epsilon$  is assumed to follow a **normal distribution**, the risk score  $\mathbf{x}^T \mathbf{w}$  is also called the **probit**.



**Figure:** pdf of the normal distribution

For our purposes, there is not much difference between *logistic* and *logit* regression. The main difference is that logistic regression is numerically easier to solve.

# Logistic Model

From now on, we'll only look at the logistic model. Note that similar derivations could be made for any other model.

Consider  $p(y = 1|\mathbf{x})$ , the **likelihood** that the output is a success:

$$\begin{aligned} p(y = 1|\mathbf{x}) &= p(\mathbf{x}^\top \mathbf{w} + \epsilon > 0) \\ &= p(\epsilon > -\mathbf{x}^\top \mathbf{w}) \end{aligned}$$

since  $\epsilon$  is symmetrically distributed around 0, it follows that


$$p(y = 1|\mathbf{x}) = p(\epsilon < \mathbf{x}^\top \mathbf{w})$$

Because we have made some assumptions about the distribution of  $\epsilon$ , we are able to derive a closed-form expression for the likelihood.

# The Logistic Function

The function  $f : t \mapsto f(t) = p(\epsilon < t)$  is the c.d.f. of the logistic distribution and is also called the **logistic function** or **sigmoid**:

$$f(t) = \frac{1}{1 + e^{-t}}$$



Thus we have a simple model for the likelihood of success  $h_{\mathbf{w}}(\mathbf{x}) = p(y = 1|\mathbf{x})$ :

$$h_{\mathbf{w}}(\mathbf{x}) = p(y = 1|\mathbf{x}) = p(\epsilon < \mathbf{x}^T \mathbf{w}) = f(\mathbf{x}^T \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{x}^T \mathbf{w}}}$$

The likelihood of failure is simply given by:

$$p(y = 0|\mathbf{x}) = 1 - h_{\mathbf{w}}(\mathbf{x})$$

**Exercise:**

show that  $p(y = 0|\mathbf{x}) = h_{\mathbf{w}}(-\mathbf{x})$

In **linear regression**, the model  $h_{\mathbf{w}}(\mathbf{x})$  was a direct prediction of the outcome:

$$h_{\mathbf{w}}(\mathbf{x}) = y$$

In **logistic regression**, the model  $h_{\mathbf{w}}(\mathbf{x})$  makes an estimation of the **likelihood** of the outcome:

$$h_{\mathbf{w}}(\mathbf{x}) = p(y = 1|\mathbf{x})$$

Thus whereas in linear regression we try to answer the question:

*What is the value of  $y$  given  $\mathbf{x}$ ?*


In logistic regression (and any other general linear model), we try instead to answer the question:

*What is the probability that  $y = 1$  given  $\mathbf{x}$ ?*



# Logistic Regression

Below is the plot of an estimated  $h_{\mathbf{w}}(\mathbf{x}) \approx p(y = 1|\mathbf{x})$  for our problem:



The results are easy to interpret: there is about 60% chance to pass the exam if you study for 3 hours.

## Maximum Likelihood

To estimate the weights  $\mathbf{w}$ , we will again use the concept of **Maximum Likelihood**.

## Maximum Likelihood

As we've just seen, for a particular observation  $\mathbf{x}_i$ , the likelihood is given by:

$$p(y = y_i | \mathbf{x}_i) = \begin{cases} p(y = 1 | \mathbf{x}_i) = h_{\mathbf{w}}(\mathbf{x}_i) & \text{if } y_i = 1 \\ p(y = 0 | \mathbf{x}_i) = 1 - h_{\mathbf{w}}(\mathbf{x}_i) & \text{if } y_i = 0 \end{cases}$$

As  $y_i \in \{0, 1\}$ , this can be rewritten in a slightly more compact form as:

$$p(y = y_i | \mathbf{x}_i) = h_{\mathbf{w}}(\mathbf{x}_i)^{y_i} (1 - h_{\mathbf{w}}(\mathbf{x}_i))^{1-y_i}$$

This works because  $z^0 = 1$ .

The likelihood over all observations is then:

$$p(\mathbf{y} | \mathbf{X}) = \prod_{i=1}^n h_{\mathbf{w}}(\mathbf{x}_i)^{y_i} (1 - h_{\mathbf{w}}(\mathbf{x}_i))^{1-y_i}$$

## Maximum Likelihood

We want to find  $\mathbf{w}$  that maximises the likelihood  $p(\mathbf{y}|\mathbf{X})$ . As always, it is equivalent but more convenient to minimise the negative log likelihood:

$$\begin{aligned} E(\mathbf{w}) &= -\ln(p(\mathbf{y}|\mathbf{X})) \\ &= \sum_{i=1}^n -y_i \ln(h_{\mathbf{w}}(\mathbf{x}_i)) - (1 - y_i) \ln(1 - h_{\mathbf{w}}(\mathbf{x}_i)) \end{aligned}$$

This error function we need to minimise is called the **cross-entropy**.

In the Machine Learning community, the error function is also frequently called a **loss function**. Thus here we would say: the loss function is the cross-entropy.

We could have considered optimising the parameters  $\mathbf{w}$  using other error functions. For instance we could have tried to minimise the least square error as we did in linear regression:

$$E_{LS}(\mathbf{w}) = \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$$

The problem is that  $h_{\mathbf{w}}$  is non-convex and makes the minimisation of  $E_{LS}(\mathbf{w})$  much harder than when using cross-entropy.

In fact, this is a mistake that the Neural Net community did for a number of years before switching to the cross entropy loss function.

## Optimisation: gradient descent

To minimise the error function, we need to resort to **gradient descent**, which is a general method for nonlinear optimisation and which will be at the core of neural networks optimisation.

We start at  $\mathbf{w}^{(0)}$  and take steps along the steepest direction  $\mathbf{v}$  using a fixed size step as follows:

$$\mathbf{w}^{(n+1)} = \mathbf{w}^{(n)} + \eta \mathbf{v}^{(n)}$$

$\eta$  is called the **learning rate** and controls the speed of the descent.

What is the steepest slope  $\mathbf{v}$ ?

## Optimisation: gradient descent

Without loss of generality we set  $\mathbf{v}$  to be a unit vector (ie.  $\|\mathbf{v}\| = 1$ ). Then, moving  $\mathbf{w}$  to  $\mathbf{w} + \eta\mathbf{v}$  yields a new error as follows:

$$E(\mathbf{w} + \eta\mathbf{v}) = E(\mathbf{w}) + \eta \left( \frac{\partial E}{\partial \mathbf{w}} \right)^{\top} \mathbf{v} + O(\eta^2)$$

which reaches a minimum when

$$\mathbf{v} = - \frac{\frac{\partial E}{\partial \mathbf{w}}}{\left\| \frac{\partial E}{\partial \mathbf{w}} \right\|}$$

## Optimisation: gradient descent

now, it is hard to set the right size for a fixed learning rate  $\eta$ .

Thus, instead of using

$$\mathbf{w}^{(n+1)} = \mathbf{w}^{(n)} - \eta \frac{\frac{\partial E}{\partial \mathbf{w}}}{\left\| \frac{\partial E}{\partial \mathbf{w}} \right\|}$$

we repeat the following update step:

$$\mathbf{w}^{(n+1)} = \mathbf{w}^{(n)} - \eta \frac{\partial E}{\partial \mathbf{w}}$$



## Optimisation: gradient descent

Recall that the cross-entropy error function to minimise is:

$$E = \sum_{i=1}^n -y_i \ln(h_{\mathbf{w}}(\mathbf{x}_i)) - (1 - y_i) \ln(1 - h_{\mathbf{w}}(\mathbf{x}_i))$$

and that  $h_{\mathbf{w}}(\mathbf{x}) = f(\mathbf{x}^T \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{x}^T \mathbf{w}}}$

### *Exercise:*

Given that the derivative of the sigmoid  $f$  is  $f'(t) = (1 - f(t))f(t)$ , show that

$$\frac{\partial E}{\partial \mathbf{w}} = \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \mathbf{x}_i$$

## Optimisation: gradient descent

The overall gradient descent method looks like so:


1. set an initial weight vector  $\mathbf{w}^{(0)}$  and
2. **for**  $t = 0, 1, 2, \dots$  **do until convergence**
3.     compute the gradient

$$\frac{\partial E}{\partial \mathbf{w}} = \sum_{i=1}^n \left( \frac{1}{1 + e^{-\mathbf{x}_i^T \mathbf{w}}} - y_i \right) \mathbf{x}_i$$

4.     update the weights:  $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \frac{\partial E}{\partial \mathbf{w}}$

## Example

Below is an example with 2 features.




## Example

The estimate for the probability of success is

$$h_{\mathbf{w}}(\mathbf{x}) = 1/(1 + e^{-(-1.28 - 1.09x_1 + 1.89x_2)})$$

Below are drawn the lines that correspond to  $h_{\mathbf{w}}(\mathbf{x}) = 0.05$ ,  $h_{\mathbf{w}}(\mathbf{x}) = 0.5$  and  $h_{\mathbf{w}}(\mathbf{x}) = 0.95$ .



# Multiclass Classification

---

It is very often that you have to deal with more than 2 classes.

## one-vs.-all

The simplest way to consider a problem that has more than 2 classes is to adopt the **one-vs.-all** (or one-against-all) strategy:

For each class  $k$ , you can train a single binary classifier ( $y = 0$  for all other class, and  $y = 1$  for class  $k$ ). The classifiers return a real-valued likelihood for their decision.

The one-vs.-all prediction returns the label for which the corresponding classifier reports the highest likelihood.

The **one-vs.-all** approach is a very simple one. However it is an heuristic that has many problems.

One problem is that for each binary classifier, the negative samples (from all the classes but  $k$ ) are more numerous and more heterogeneous than the positive samples (from class  $k$ ).

A better approach would be to have a unified model for all classifiers and jointly train them. The extension of Logistic regression that just does this is called **multinomial logistic regression**.



## Multinomial Logistic Regression

In Multinomial Logistic Regression, each of the binary classifier is based on the following likelihood model:

$$p(y = C_k | \mathbf{x}) = \text{softmax}(\mathbf{x}^\top \mathbf{w})_k = \frac{\exp(\mathbf{w}_k^\top \mathbf{x})}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \mathbf{x})}$$

$C_k$  is the class  $k$  and  $\text{softmax} : \mathbb{R}^K \rightarrow \mathbb{R}^K$  is the function defined as

$$\text{softmax}(\mathbf{t})_k = \frac{\exp(t_k)}{\sum_{j=1}^K \exp(t_j)}$$

In other words, **softmax** takes as an input the vector of logits for all classes and returns the vector of corresponding likelihoods.

## Softmax Optimisation

To optimise for the parameters. We can take again the **maximum likelihood** approach.

Combining the likelihood for all possible classes gives us:

$$p(y|\mathbf{x}) = p(y = C_1|\mathbf{x})^{[y=C_1]} \times \dots \times p(y = C_K|\mathbf{x})^{[y=C_K]}$$

where again  $[y = C_1]$  is 1 if  $y = C_1$  and 0 otherwise.

The total likelihood is:

$$p(y|\mathbf{X}) = \prod_{i=1}^n p(y_i = C_1|\mathbf{x}_i)^{[y=C_1]} \times \dots \times p(y_i = C_K|\mathbf{x}_i)^{[y=C_K]}$$

Taking the negative log likelihood yields the cross entropy error function for the multiclass problem:

$$E(\mathbf{w}_1, \dots, \mathbf{w}_K) = -\ln(p(y|\mathbf{X})) = -\sum_{i=1}^n \sum_{k=1}^K [y_i = C_k] \ln(p(y_i = C_k | \mathbf{x}_i))$$

Similarly to logistic regression, we can use a gradient descent approach to find the  $K$  weight vectors  $\mathbf{w}_1, \dots, \mathbf{w}_K$  that minimise this cross entropy expression.

## Take Away

With **Logistic Regression**, we look at linear models, where the output of the problem is a **binary categorical** response.

Instead of directly predicting the actual outcome as in least squares, the model proposed in logistic regression makes a prediction about the **likelihood of belonging to a particular class**.

Finding the maximum likelihood parameters is equivalent to minimising the **cross entropy** loss function. The minimisation can be done using the **gradient descent** technique.

The extension of Logistic Regression to more than 2 classes is called the **Multinomial Logistic Regression**.